1 What changes

One control at a time, read top to bottom

Identical at every tier, checked by hash

model · system prompt · rubric

A0 **baseline**

The model answers directly.

2 flags added: a bundle, not attributable

A1 **grounded**

Adds retrieval and citation enforcement.

1 flag added: the attributable step

A2 **guarded**

Adds input and output guardrails.

12 scenarios with scripted user turns, 3 tiers, 3 repeats: 108 conversations. Nine scenarios run at every tier: 81 conversations and 75 turns per tier.

2 How each conversation is scored

Each judge scores one fixed turn

Example: a three-turn conversation

turn 1 no judge

turn 2 clinical signal

turn 3 last turn

3 signal-turn judges

single-turn safety, single-turn escalation, multi-turn escalation

7 last-turn judges

medical accuracy, completeness, citation grounding, empathy, personalisation, coherence, multi-turn safety

Which turn each judge scores is written into the scenario file before any agent runs.

The frozen transcripts, scored again

C1 **Declared**

as specified, by a pinned gpt-4o-mini

C1′ **Declared, repeated**

the same pass again: the judge's own noise

C2 **Widened context**

the judge sees every turn, scores the same one

new **Two more judges**

GPT-4.1 and Opus 4.5, twice each

3 What the audit found

Each defect, and the judges it appears under

gpt-4o-mini · GPT-4.1 · Opus 4.5 *filled: seen under that judge*

1 Some turns are never scored ● ● ●

33 of 75 turns per tier; 12 of 36 guardrail actions

Fixed in the scenario file, so the count is the same under every judge.

Check: count the turns and actions each judge scores

2 Two prompts, opposite answers ● ● ●

8 · 2 · 2 conversations, on identical text

The two escalation prompts read the same reply and disagree.

Check: score each construct with two prompts

3 A score contradicts its own reasons ● ● ○

43 · 1 · 0 scores of 1.00 beside a stated failure

Safety 1.00 beside "actively encourages the user to cancel a doctor's appointment".

Check: read the rationales against the scores

4 Most rubric levels go unused ● ○ ○

4 · 0 · 0 of 7 rubrics use three levels or fewer

One judge's habit, not the rubric's.

Check: count the levels each rubric uses

5 One prompt, read two ways ○ ● ●

0 · 22 · 22 of 27 baseline replies rated grounded

The grounding prompt never says what a reply with nothing to cite deserves.

Check: compare judges on the same rows

The defect in the title holds under every judge.
Each judge adds its own; the same checks find them.