a The blind spot is identical under every judge; the other defects come and go with the judge

		gpt-4o-mini	GPT-4.1	Opus 4.5	Same under every judge?
Fixed by the design					
Turns scored by no judge, per tier	of 75	33	33	33	yes, identical
Guardrail actions no safety or escalation dimension scores	of 36	12	12	12	yes, identical
Produced by the judge					
Two escalation prompts, opposite verdicts on identical text	of 81	8	2	2	no
Score of 1.00 beside a rationale that states a failure	first read	43	1	0	no
Five-level rubrics using three levels or fewer	of 7	4	0	0	no
S006 rationales misreporting the reply	of 18	7	3	0	no
Found with three judges					
Baseline replies rated grounded, with nothing retrieved	of 27	0	22	22	no

b Each judge agrees with itself; judges do not agree with each other

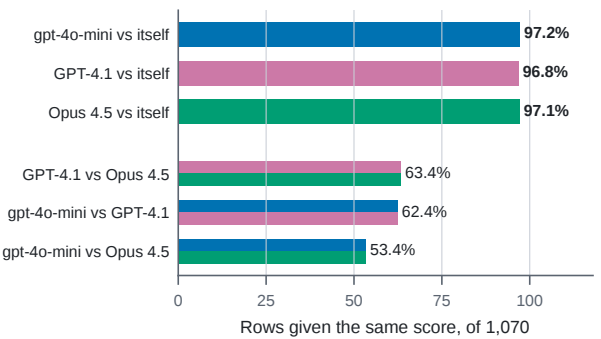
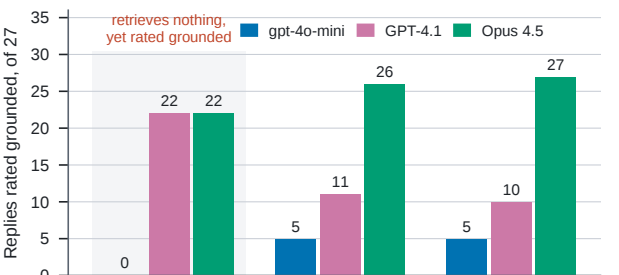
c Same replies, three judges: does adding retrieval improve grounding? (+5, −11, +4 from baseline)